

yo-bro

July 12, 2024

```
[1]: import os
import numpy as np
import matplotlib.pyplot as plt
import warnings
from tqdm.notebook import tqdm
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator,
    array_to_img
from tensorflow.keras.models import Sequential
from tensorflow.keras import layers
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import BinaryCrossentropy

warnings.filterwarnings('ignore')
```

```
[3]: BASE_DIR = 'E:\\cbnslab\\GAN_proj\\images'
TARGET_SIZE = (64, 64) # Set appropriate target size
BATCH_SIZE = 32
```

```
[5]: # Create an ImageDataGenerator instance
datagen = ImageDataGenerator(rescale=1./127.5, preprocessing_function=lambda x:
    (x - 1))
```

```
[13]: # Load data from the directory
dataset = tf.keras.utils.image_dataset_from_directory(
    BASE_DIR,
    label_mode=None, # No labels
    image_size=TARGET_SIZE,
    batch_size=BATCH_SIZE,
    shuffle=True
)
```

Found 21527 files.

```
[15]: # Normalize the images
normalization_layer = tf.keras.layers.Rescaling(1./127.5, offset=-1)
```

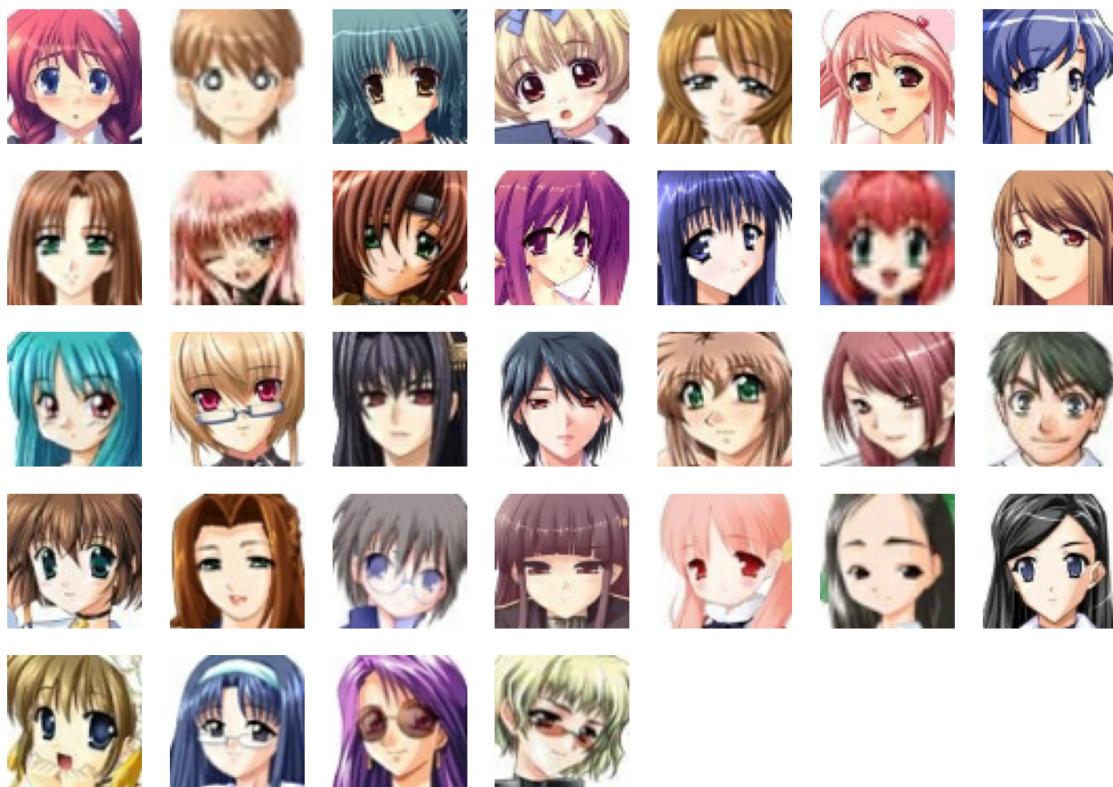
```

dataset = dataset.map(lambda x: normalization_layer(x))

[17]: # Get a batch of images
for batch in dataset.take(1):
    sample_images = batch.numpy()

plt.figure(figsize=(20, 20))
for i in range(min(len(sample_images), 49)):
    plt.subplot(7, 7, i+1)
    plt.imshow((sample_images[i] + 1) * 127.5 / 255)
    plt.axis('off')
plt.show()

```



```

[19]: LATENT_DIM = 100
WEIGHT_INIT = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)

# Generator
def build_generator():
    model = Sequential(name='generator')
    model.add(layers.Dense(8 * 8 * 512, input_dim=LATENT_DIM))
    model.add(layers.ReLU())
    model.add(layers.Reshape((8, 8, 512)))

```

```

    model.add(layers.Conv2DTranspose(256, (4, 4), strides=(2, 2),  

↳padding='same', kernel_initializer=WEIGHT_INIT))  

    model.add(layers.ReLU())  

    model.add(layers.Conv2DTranspose(128, (4, 4), strides=(2, 2),  

↳padding='same', kernel_initializer=WEIGHT_INIT))  

    model.add(layers.ReLU())  

    model.add(layers.Conv2DTranspose(64, (4, 4), strides=(2, 2),  

↳padding='same', kernel_initializer=WEIGHT_INIT))  

    model.add(layers.ReLU())  

    model.add(layers.Conv2D(3, (4, 4), padding='same', activation='tanh'))  

    return model

# Discriminator
def build_discriminator(input_shape):  

    model = Sequential(name='discriminator')  

    model.add(layers.Conv2D(64, (4, 4), strides=(2, 2), padding='same',  

↳input_shape=input_shape))  

    model.add(layers.LeakyReLU(alpha=0.2))  

    model.add(layers.Conv2D(128, (4, 4), strides=(2, 2), padding='same'))  

    model.add(layers.LeakyReLU(alpha=0.2))  

    model.add(layers.Conv2D(256, (4, 4), strides=(2, 2), padding='same'))  

    model.add(layers.LeakyReLU(alpha=0.2))  

    model.add(layers.Flatten())  

    model.add(layers.Dropout(0.3))  

    model.add(layers.Dense(1, activation='sigmoid'))  

    return model

generator = build_generator()  

discriminator = build_discriminator((64, 64, 3))

generator.summary()  

discriminator.summary()

```

Model: "generator"

Layer (type)	Output Shape	
Param #		
dense (Dense) ↳ 3,309,568	(None, 32768)	
re_lu (ReLU) ↳ 0	(None, 32768)	

reshape (Reshape)	(None, 8, 8, 512)	□
↳ 0		
conv2d_transpose (Conv2DTranspose)	(None, 16, 16, 256)	□
↳ 2,097,408		
re_lu_1 (ReLU)	(None, 16, 16, 256)	□
↳ 0		
conv2d_transpose_1 (Conv2DTranspose)	(None, 32, 32, 128)	□
↳ 524,416		
re_lu_2 (ReLU)	(None, 32, 32, 128)	□
↳ 0		
conv2d_transpose_2 (Conv2DTranspose)	(None, 64, 64, 64)	□
↳ 131,136		
re_lu_3 (ReLU)	(None, 64, 64, 64)	□
↳ 0		
conv2d (Conv2D)	(None, 64, 64, 3)	□
↳ 3,075		

Total params: 6,065,603 (23.14 MB)

Trainable params: 6,065,603 (23.14 MB)

Non-trainable params: 0 (0.00 B)

Model: "discriminator"

Layer (type)	Output Shape	□
Param #		
conv2d_1 (Conv2D)	(None, 32, 32, 64)	□
↳ 3,136		
leaky_re_lu (LeakyReLU)	(None, 32, 32, 64)	□
↳ 0		

conv2d_2 (Conv2D)	(None, 16, 16, 128)	□
↪ 131,200		
leaky_re_lu_1 (LeakyReLU)	(None, 16, 16, 128)	□
↪ 0		
conv2d_3 (Conv2D)	(None, 8, 8, 256)	□
↪ 524,544		
leaky_re_lu_2 (LeakyReLU)	(None, 8, 8, 256)	□
↪ 0		
flatten (Flatten)	(None, 16384)	□
↪ 0		
dropout (Dropout)	(None, 16384)	□
↪ 0		
dense_1 (Dense)	(None, 1)	□
↪ 16,385		

Total params: 675,265 (2.58 MB)

Trainable params: 675,265 (2.58 MB)

Non-trainable params: 0 (0.00 B)

```
[21]: class DCGAN(keras.Model):
    def __init__(self, generator, discriminator, latent_dim):
        super().__init__()
        self.generator = generator
        self.discriminator = discriminator
        self.latent_dim = latent_dim
        self.g_loss_metric = keras.metrics.Mean(name='g_loss')
        self.d_loss_metric = keras.metrics.Mean(name='d_loss')

    @property
    def metrics(self):
        return [self.g_loss_metric, self.d_loss_metric]

    def compile(self, g_optimizer, d_optimizer, loss_fn):
        super(DCGAN, self).compile()
        self.g_optimizer = g_optimizer
```

```

    self.d_optimizer = d_optimizer
    self.loss_fn = loss_fn

    def train_step(self, real_images):
        batch_size = tf.shape(real_images)[0]
        random_noise = tf.random.normal(shape=(batch_size, self.latent_dim))

        with tf.GradientTape() as tape:
            pred_real = self.discriminator(real_images, training=True)
            real_labels = tf.ones((batch_size, 1)) * 0.9
            d_loss_real = self.loss_fn(real_labels, pred_real)

            fake_images = self.generator(random_noise)
            pred_fake = self.discriminator(fake_images, training=True)
            fake_labels = tf.zeros((batch_size, 1))
            d_loss_fake = self.loss_fn(fake_labels, pred_fake)

            d_loss = (d_loss_real + d_loss_fake) / 2

            gradients = tape.gradient(d_loss, self.discriminator.
                                      trainable_variables)
            self.d_optimizer.apply_gradients(zip(gradients, self.discriminator.
                                      trainable_variables))

        with tf.GradientTape() as tape:
            fake_images = self.generator(random_noise, training=True)
            pred_fake = self.discriminator(fake_images, training=True)
            g_loss = self.loss_fn(tf.ones((batch_size, 1)), pred_fake)

            gradients = tape.gradient(g_loss, self.generator.trainable_variables)
            self.g_optimizer.apply_gradients(zip(gradients, self.generator.
                                      trainable_variables))

            self.d_loss_metric.update_state(d_loss)
            self.g_loss_metric.update_state(g_loss)

        return {'d_loss': self.d_loss_metric.result(), 'g_loss': self.
                g_loss_metric.result()}

```

```
[25]: class DCGANMonitor(keras.callbacks.Callback):
    def __init__(self, num_imgs=25, latent_dim=100):
        self.num_imgs = num_imgs
        self.latent_dim = latent_dim
        self.noise = tf.random.normal([25, latent_dim])

    def on_epoch_end(self, epoch, logs=None):
        g_img = self.model.generator(self.noise)
```

```

g_img = (g_img * 127.5) + 127.5
fig = plt.figure(figsize=(8, 8))
for i in range(self.num_imgs):
    plt.subplot(5, 5, i+1)
    img = array_to_img(g_img[i])
    plt.imshow(img)
    plt.axis('off')
plt.show()

def on_train_end(self, logs=None):
    self.model.generator.save('generator.h5')

dcgan = DCGAN(generator=generator, discriminator=discriminator,
               latent_dim=LATENT_DIM)

D_LR = 0.0001
G_LR = 0.0003
dcgan.compile(g_optimizer=Adam(learning_rate=G_LR, beta_1=0.5),
               d_optimizer=Adam(learning_rate=D_LR, beta_1=0.5),
               loss_fn=BinaryCrossentropy())

N_EPOCHS = 25
dcgan.fit(dataset, epochs=N_EPOCHS, callbacks=[DCGANMonitor()])

```

Epoch 1/25  
673/673 0s 821ms/step -  
d\_loss: 0.6929 - g\_loss: 0.8388



673/673 558s 822ms/step -

d\_loss: 0.6929 - g\_loss: 0.8388

Epoch 2/25

673/673 0s 872ms/step -

d\_loss: 0.6891 - g\_loss: 0.8621



673/673 587s 873ms/step -

d\_loss: 0.6891 - g\_loss: 0.8621

Epoch 3/25

673/673 0s 804ms/step -

d\_loss: 0.6883 - g\_loss: 0.8743



673/673 542s 805ms/step -

d\_loss: 0.6883 - g\_loss: 0.8743

Epoch 4/25

673/673 0s 816ms/step -

d\_loss: 0.6893 - g\_loss: 0.8570



673/673 550s 817ms/step -

d\_loss: 0.6893 - g\_loss: 0.8570

Epoch 5/25

673/673 0s 802ms/step -

d\_loss: 0.6845 - g\_loss: 0.8622



673/673 542s 806ms/step -

d\_loss: 0.6845 - g\_loss: 0.8622

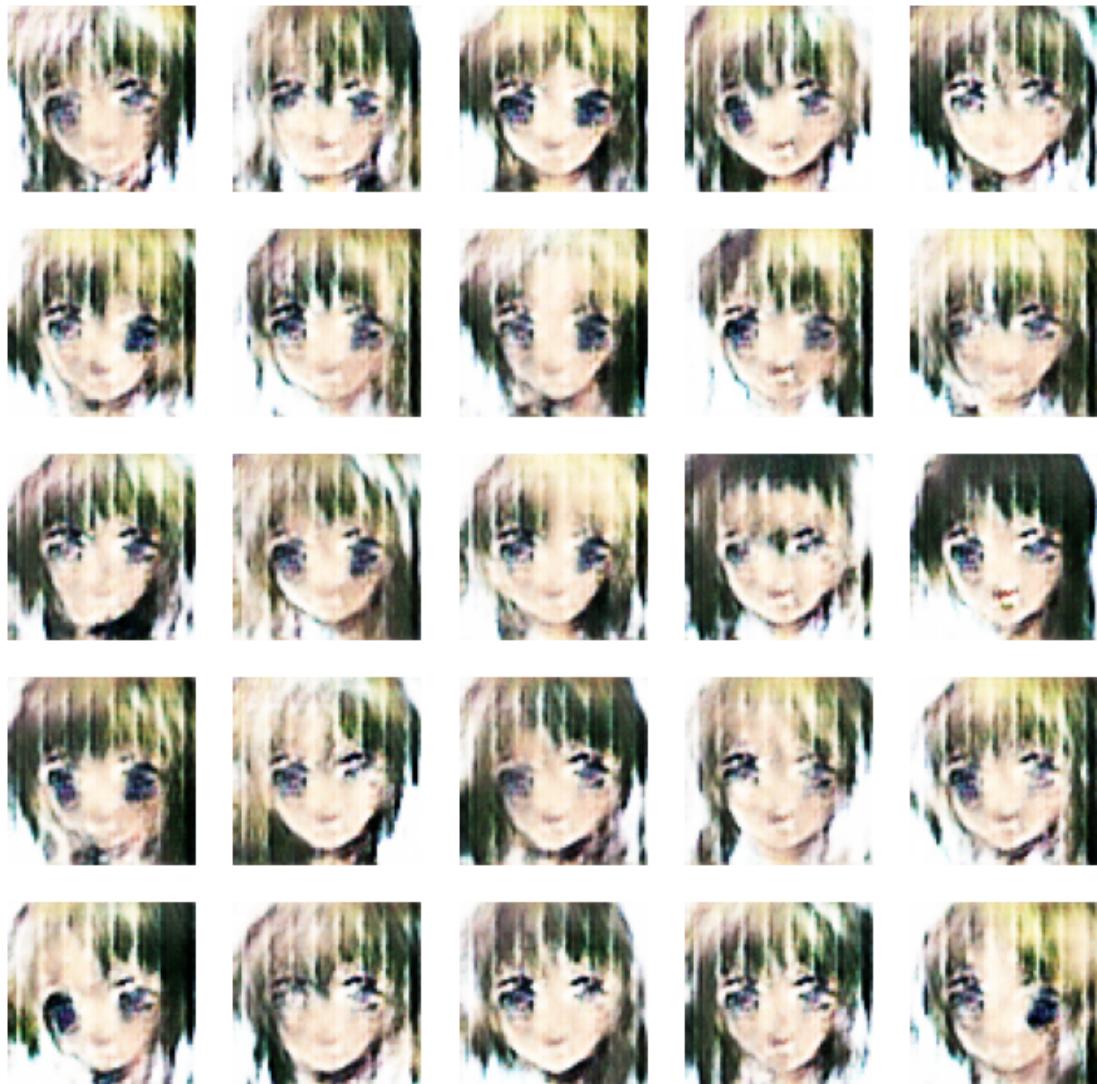
Epoch 6/25

673/673 0s 801ms/step -

d\_loss: 0.6789 - g\_loss: 0.8996



673/673 539s 801ms/step -  
d\_loss: 0.6789 - g\_loss: 0.8996  
Epoch 7/25  
673/673 0s 791ms/step -  
d\_loss: 0.6827 - g\_loss: 0.9067



673/673 533s 792ms/step -

d\_loss: 0.6827 - g\_loss: 0.9067

Epoch 8/25

673/673 0s 775ms/step -

d\_loss: 0.6800 - g\_loss: 0.8994



673/673 522s 776ms/step -

d\_loss: 0.6800 - g\_loss: 0.8994

Epoch 9/25

673/673 0s 783ms/step -

d\_loss: 0.6750 - g\_loss: 0.9359



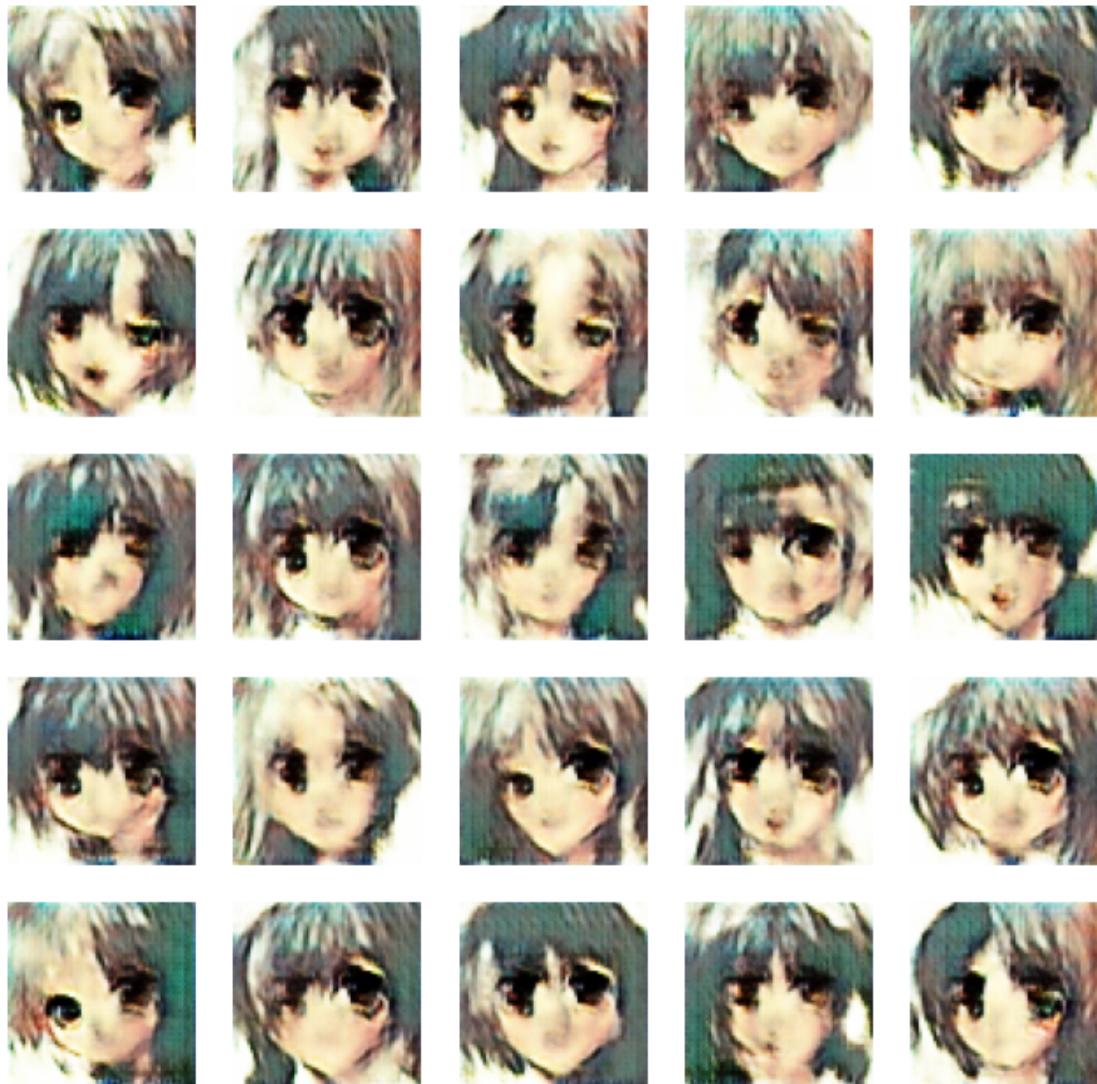
673/673 528s 784ms/step -

d\_loss: 0.6750 - g\_loss: 0.9359

Epoch 10/25

673/673 0s 810ms/step -

d\_loss: 0.6624 - g\_loss: 0.9789



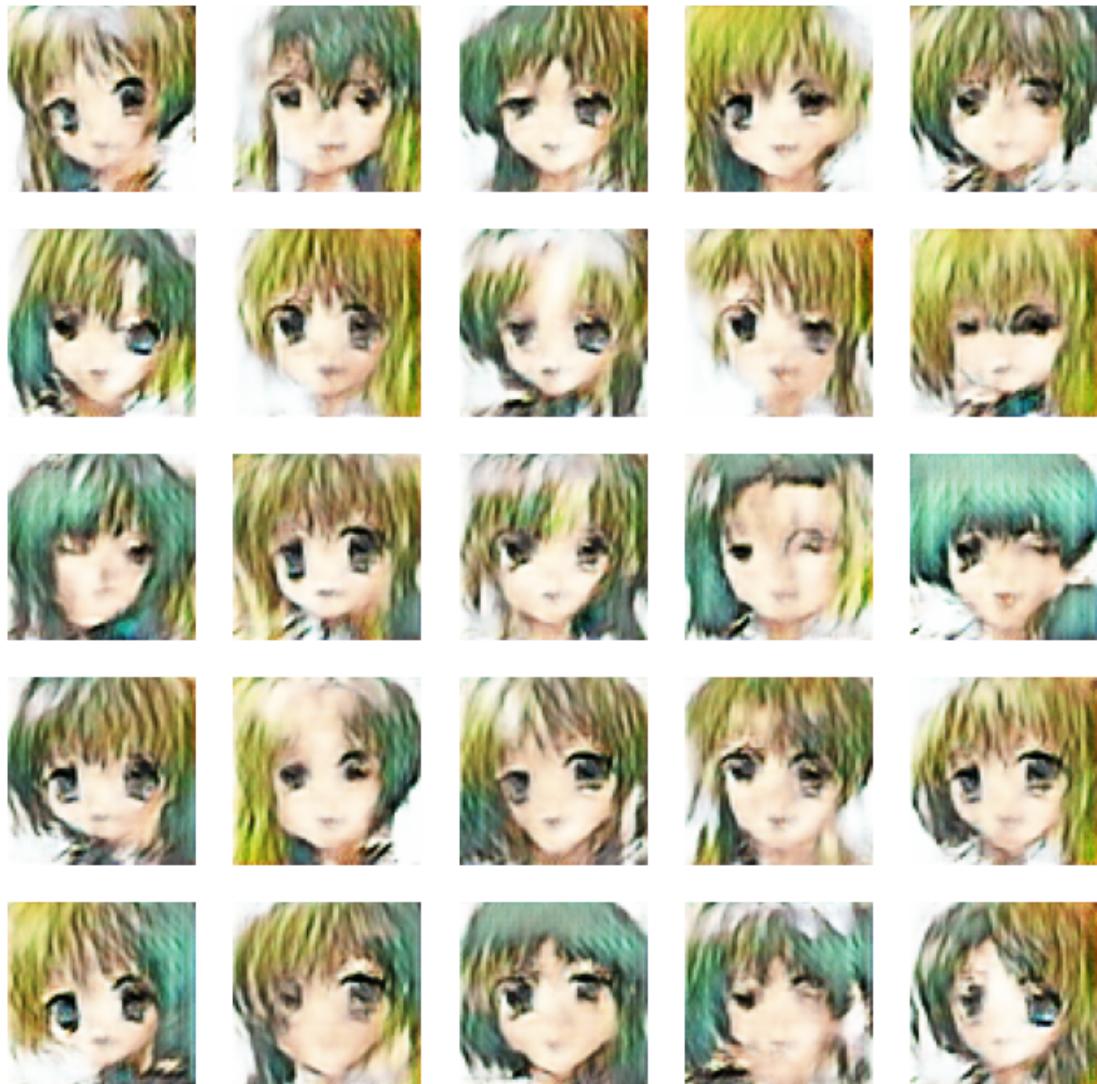
673/673 546s 812ms/step -

d\_loss: 0.6624 - g\_loss: 0.9789

Epoch 11/25

673/673 0s 808ms/step -

d\_loss: 0.6664 - g\_loss: 0.9581



673/673 545s 809ms/step -

d\_loss: 0.6664 - g\_loss: 0.9581

Epoch 12/25

673/673 0s 816ms/step -

d\_loss: 0.6717 - g\_loss: 0.9402



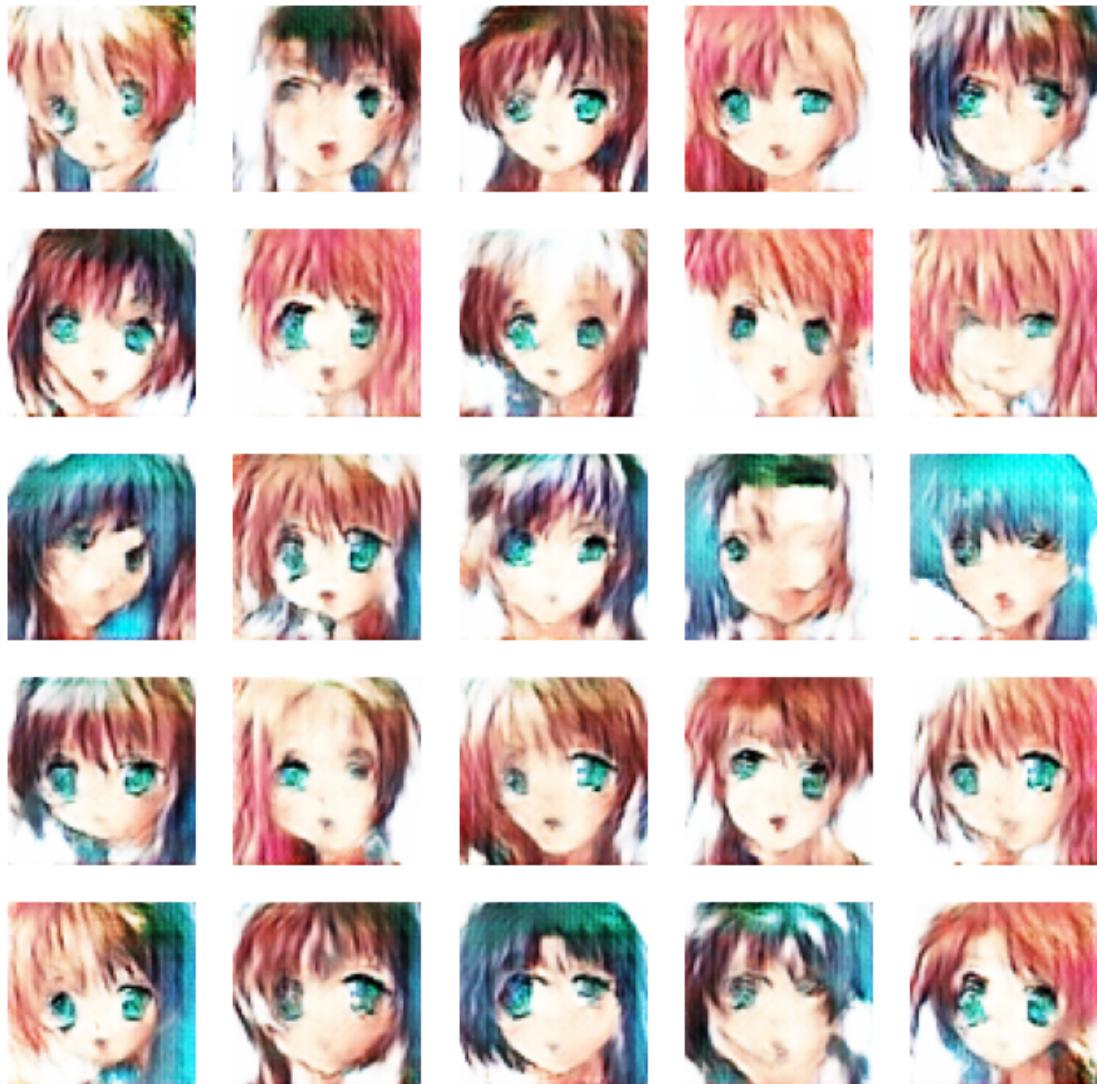
673/673 550s 818ms/step -

d\_loss: 0.6717 - g\_loss: 0.9402

Epoch 13/25

673/673 0s 810ms/step -

d\_loss: 0.6753 - g\_loss: 0.9134



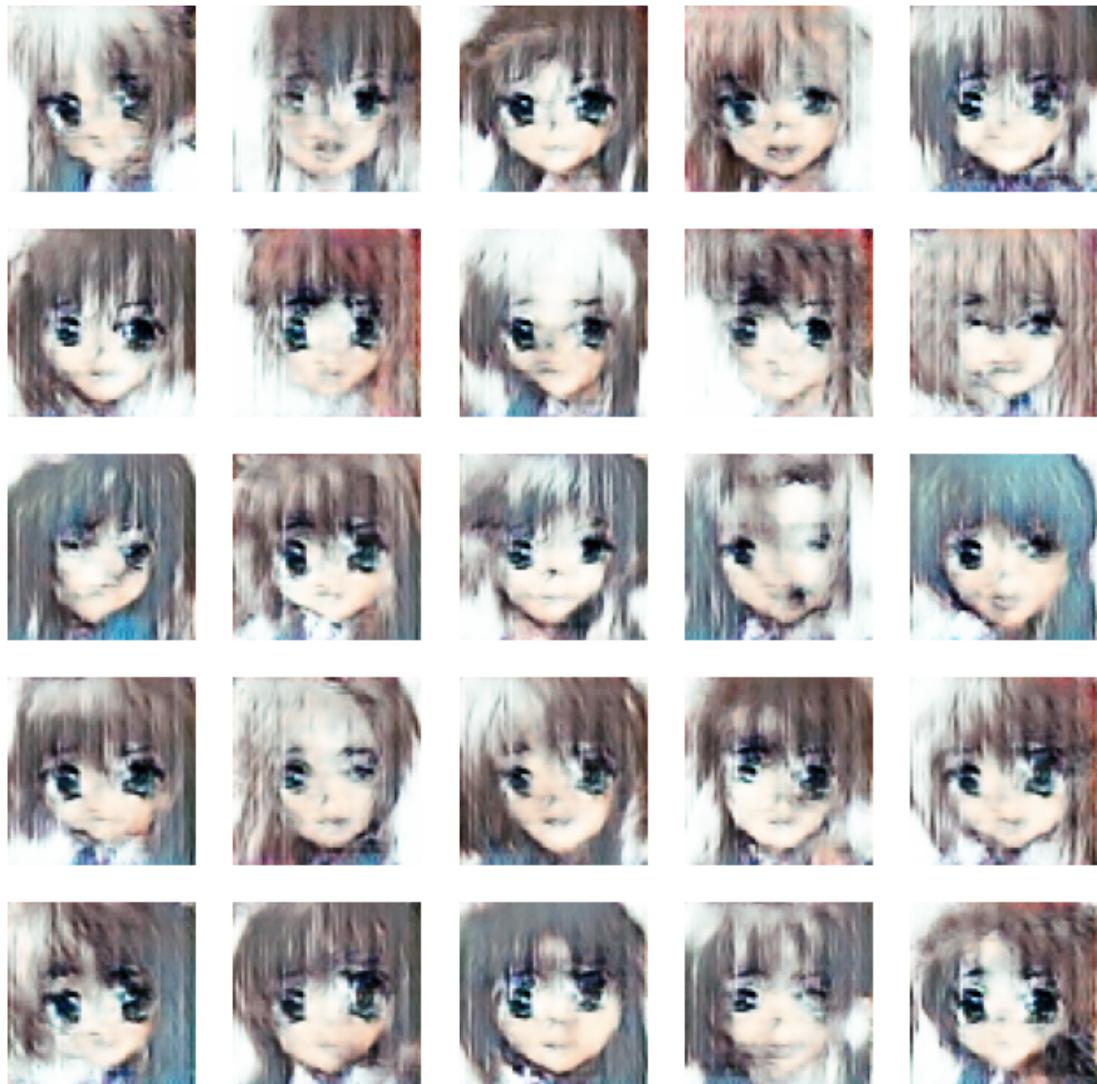
673/673 546s 811ms/step -

d\_loss: 0.6753 - g\_loss: 0.9134

Epoch 14/25

673/673 0s 809ms/step -

d\_loss: 0.6761 - g\_loss: 0.9063



673/673 547s 813ms/step -

d\_loss: 0.6761 - g\_loss: 0.9063

Epoch 15/25

673/673 0s 809ms/step -

d\_loss: 0.6673 - g\_loss: 0.9376



673/673 546s 810ms/step -

d\_loss: 0.6673 - g\_loss: 0.9376

Epoch 16/25

673/673 0s 796ms/step -

d\_loss: 0.6683 - g\_loss: 0.9356



673/673 537s 797ms/step -

d\_loss: 0.6683 - g\_loss: 0.9357

Epoch 17/25

673/673 0s 795ms/step -

d\_loss: 0.6656 - g\_loss: 0.9408



673/673 536s 796ms/step -

d\_loss: 0.6656 - g\_loss: 0.9408

Epoch 18/25

673/673 0s 805ms/step -

d\_loss: 0.6666 - g\_loss: 0.9380



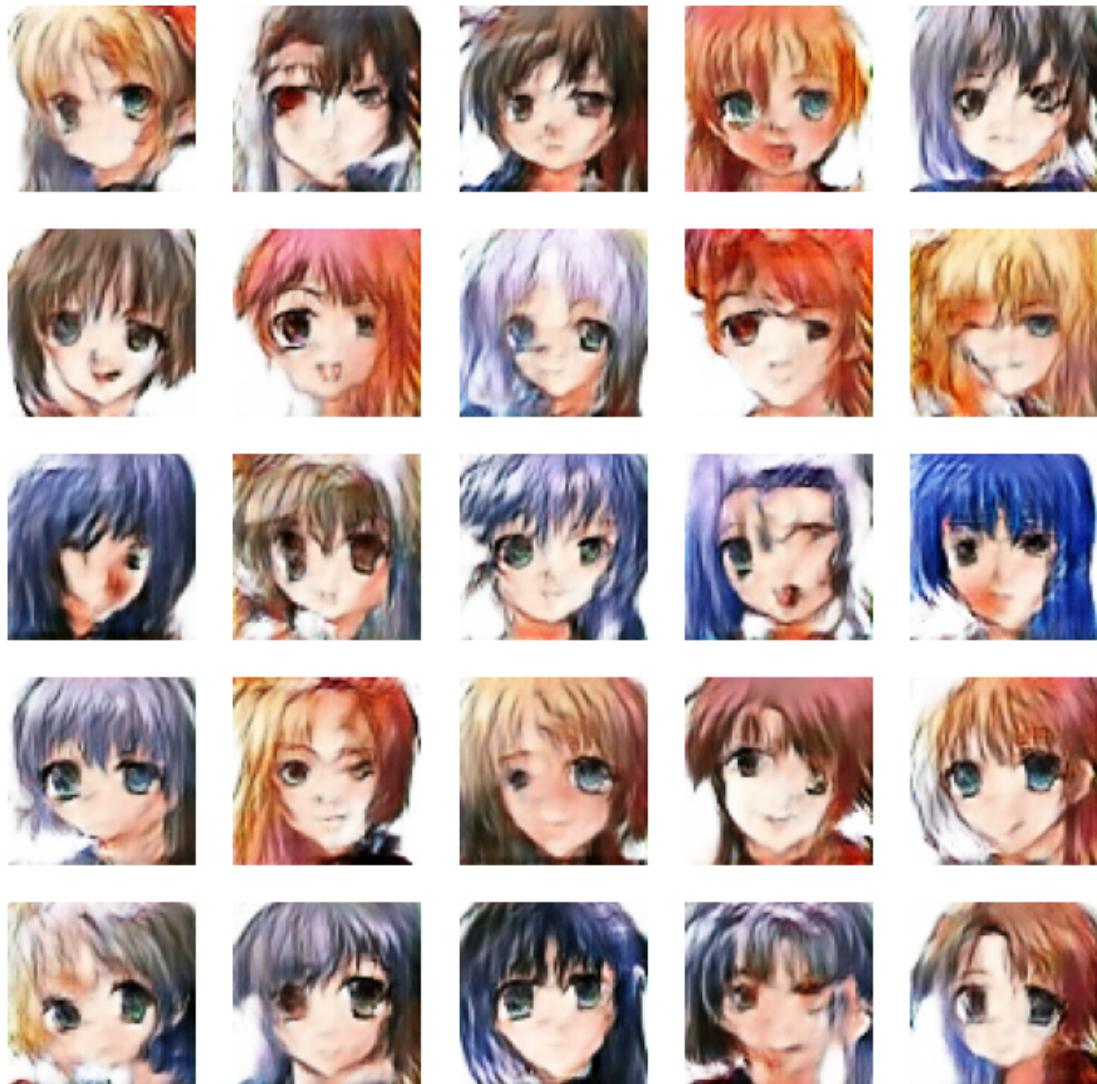
673/673 543s 806ms/step -

d\_loss: 0.6666 - g\_loss: 0.9381

Epoch 19/25

673/673 0s 816ms/step -

d\_loss: 0.6663 - g\_loss: 0.9235



673/673 551s 819ms/step -

d\_loss: 0.6663 - g\_loss: 0.9235

Epoch 20/25

673/673 0s 805ms/step -

d\_loss: 0.6674 - g\_loss: 0.9313



673/673 543s 806ms/step -

d\_loss: 0.6674 - g\_loss: 0.9313

Epoch 21/25

673/673 0s 807ms/step -

d\_loss: 0.6686 - g\_loss: 0.9288



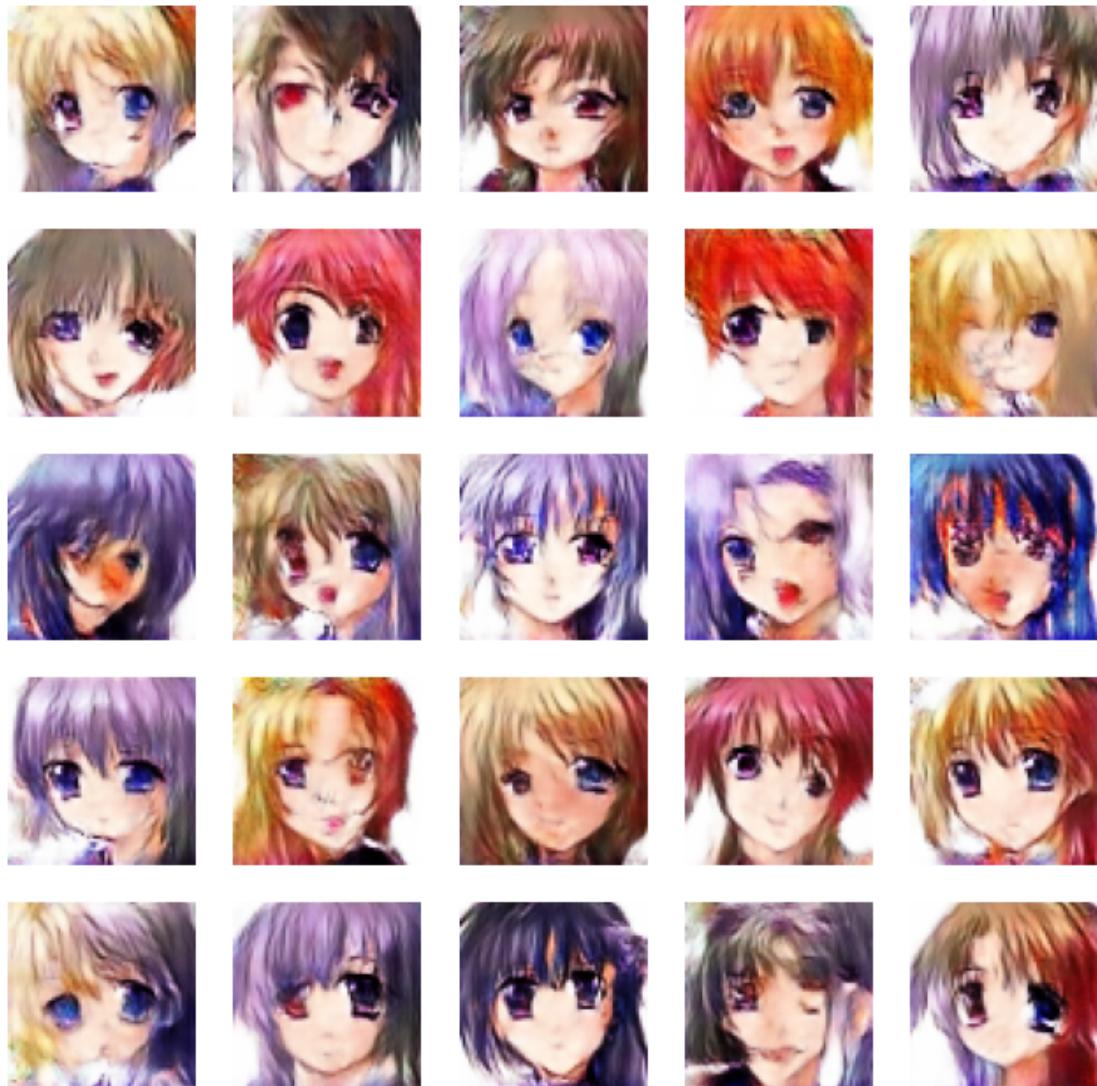
673/673 544s 808ms/step -

d\_loss: 0.6686 - g\_loss: 0.9288

Epoch 22/25

673/673 0s 813ms/step -

d\_loss: 0.6687 - g\_loss: 0.9344



673/673 548s 814ms/step -

d\_loss: 0.6687 - g\_loss: 0.9344

Epoch 23/25

673/673 0s 806ms/step -

d\_loss: 0.6687 - g\_loss: 0.9276



673/673 543s 807ms/step -

d\_loss: 0.6687 - g\_loss: 0.9276

Epoch 24/25

673/673 0s 812ms/step -

d\_loss: 0.6676 - g\_loss: 0.9338



673/673 547s 813ms/step -

d\_loss: 0.6676 - g\_loss: 0.9338

Epoch 25/25

673/673 0s 821ms/step -

d\_loss: 0.6658 - g\_loss: 0.9372



```
673/673          554s 822ms/step -  
d_loss: 0.6658 - g_loss: 0.9372
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

[25]: <keras.src.callbacks.history.History at 0x27dcb553170>

```
[31]: noise = tf.random.normal([1, 100])  
fig = plt.figure(figsize=(3, 3))  
g_img = dcgan.generator(noise)
```

```
g_img = (g_img * 127.5) + 127.5
img = array_to_img(g_img[0])
plt.imshow(img)
plt.axis('off')
plt.show()
```



[ ]: